

R: Advanced Topics

Machine Learning and Big Data

Gang Chen
chengang@genomics.cn

November 14, 2015


Outline

- 1 Overview of Machine Learning
- 2 Unsupervised Learning
- 3 Supervised Learning
- 4 Future
- 5 Big Data

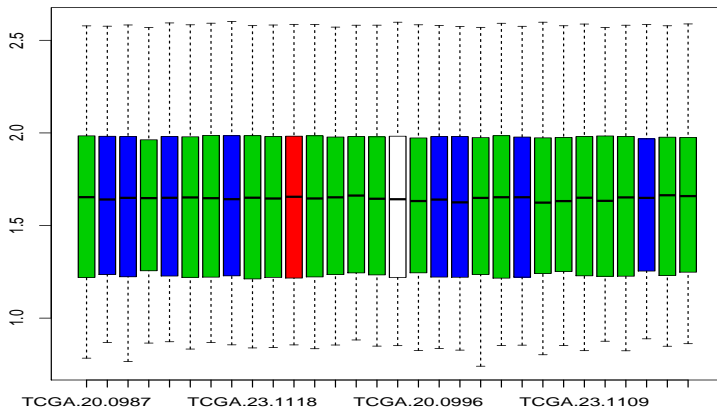
Next


- 1 Overview of Machine Learning
 - What is Data Mining?
 - Organization of the course
 - QuickStarts of R
- 2 Unsupervised Learning
- 3 Supervised Learning
- 4 Future
- 5 Big Data



Biomarker for Tumor Stage



Biomarker for Tumor Stage







How to identify tumor from genomics perspective?
Is it possible to predict tumor stage?

The answer is Data Mining.



What is Data Mining?


Pang-Ning Tan, Introduction to Data Mining

Data Mining is the process of automatically discovering useful information in large data repositories.

Knowledge Discovery in Databases

Data Mining is an integral part of knowledge discovery in databases(KDD).

Data Mining and KDD



Machine Learning and Bioinformatics

Biological Experiments

Microarray
Sequencing
Mass Spectrum
...



Preprocssing

base calling
alignment
variants
...



Data Mining

Classification
Clustering
Regression
Feature Engineering
...




Biological Knowledge

Traditional Data Analysis

Motivations

- Scalability
- High Dimensionality
- Heterogeneous and Complex Data
- Data Ownership and Distribution
- Non-traditional analysis

Data Mining Tasks



Data Mining and Machine Learning

Machine Learning

Machine learning provides the technical basis of data mining.

---Data Mining: Practical Machine Learning Tools and Techniques

Schedule

Schedule

- Introduction
- Unsupervised Learning: Clustering
- Supervised Learning: Classification
- Discussion

Softwares

Softwares

- R: R is an free platform for data analysis and visuaztion.
- R packages:
 - e1071 SVM
 - curatedOvarianData Microarry data of tumor. [Database 2013]
 - mtcars A small dataset for concept demonstration.
- Rgui
- Emacs + ESS
- Vim + R-Plugin
- RStudio
- RStudio in the Cloud


R

- Download from www.r-project.org.
- Installation and start;
- Install e1071 package.

Next

- 1 Overview of Machine Learning
- 2 Unsupervised Learning**
 - Unsupervised Learning in Bioinformatics
 - Hierarchical Clustering and its Applications in Bioinformatics
 - Summary
- 3 Supervised Learning
- 4 Future
- 5 Big Data

Heat Map



Unsupervised Learning: Clustering

What is clustering?

Cluster analysis or clustering is the task of grouping a set of objects in such a way that objects in the same group (called cluster) are more similar (in some sense or another) to each other than to those in other groups (clusters).

---Wikipedia

Applications

Applications of Clustering

- Clustering for Understanding
 - Biology
 - Information Retrieval
 - Climate
 - Psychology and medicine
 - business
 -
- Clustering for Utility
 - Summarization
 - compression
 - Efficiently finding nearest neighbors
 -

Common Clustering Methods

Clustering Methods

- Density-based clustering
- K-means
- Hierarchical Clustering
- Semi-supervised clustering
-

Unsupervised Learning in Bioinformatics

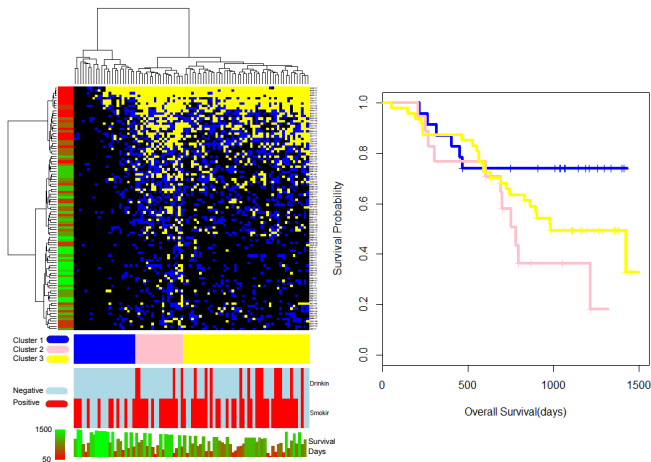
- Discovery of tumor subtypes by clustering gene expression, CNV, miRNA or integrated data.
- Clonal evolution analysis of tumor
- Mutation spectrum clustering
- Pathway or functional annotation based clustering
- Graph clustering for identification of protein functional module or protein complex
- Clustering metagenomic sequences
- Metabolomics
-

Hierarchical Clustering

Steps

- Calculating distance between individuals
- Combine closest individuals (optional, recalculate distance)
- Visualization and annotation

Clustering in Bioinformatics



Hierarchical Clustering in R

```
help(dist)  
help(hclust)  
help(heatmap)
```

Distance Calculation

```
dist(x, method = "euclidean",  
diag = FALSE,  
upper = FALSE,  
p = 2)
```

Distance Calculation

```
exprDist = dist(t(subdata))
exprDist
```

```
##          TCGA.20.0987 TCGA.23.1031 TCGA.24.0979 TCGA.23.1117
## TCGA.23.1031      6.808850
## TCGA.24.0979      6.581221      4.830934
## TCGA.23.1117      5.779257      5.328651      5.726768
## TCGA.23.1021      6.131524      4.801447      4.920307      5.456621
## TCGA.04.1337      7.030940      5.752603      6.475261      4.487964
## TCGA.20.0990      6.465878      5.308989      5.974427      5.647812
## TCGA.23.1032      7.002809      5.420443      5.256359      6.401355
## TCGA.23.1118      5.562177      5.145970      4.799568      4.599808
## TCGA.23.1026      5.388687      5.586002      5.880299      4.388108
## TCGA.20.0991      5.401875      7.061145      7.235305      5.988029
## TCGA.24.1103      6.684103      5.321956      4.508152      6.069577
## TCGA.24.0982      5.690747      4.891739      5.676629      4.789224
## TCGA.23.1119      5.727415      6.252046      6.338474      5.280068
## TCGA.23.1028      5.144723      5.505960      5.534630      5.038784
## TCGA.04.1341      6.449114      5.339794      6.528626      5.084565
## TCGA.20.0996      6.743035      5.049169      5.365110      5.371963
## TCGA.24.1104      4.451846      5.635849      5.834086      4.966425
## TCGA.23.1107      5.398379      6.024515      6.524868      5.269340
## TCGA.23.1120      5.571779      5.116315      4.304913      5.097922
## TCGA.23.1030      6.064855      4.696293      5.837826      5.243088
## TCGA.04.1342      5.574656      4.926885      5.309545      3.867411
## TCGA.23.1022      6.683327      4.765877      4.917100      5.622828
## TCGA.24.1105      5.692072      5.260009      6.253688      4.769659
## TCGA.23.1109      6.134479      5.047752      5.648282      3.975083
## TCGA.23.1121      6.502517      6.605177      7.450118      6.307436
```

Clustering

```
hclust(d, method = "complete", members = NULL)
```

Clustering

```
exprDist = dist(t(subdata))  
exprClust = hclust(exprDist)  
exprClust  
  
##  
## Call:  
## hclust(d = exprDist)  
##  
## Cluster method      : complete  
## Distance             : euclidean  
## Number of objects: 80
```


Visualization

Visualization

- Dendrogram: `plot.hclust`
- Heat Map: `heatmap`

Dendrogram

```
plot(exprClust, cex=0.6, main="Dendrogram")
```




plot.hclust

```
plot(x, labels = NULL,  
     hang = 0.1,  
     axes = TRUE,  
     frame.plot = FALSE,  
     ann = TRUE,  
     main = "Cluster Dendrogram",  
     sub = NULL,  
     xlab = NULL, ylab = "Height", ...)
```

Heat Map


```
heatmap(subdata)
```



heatmap

```
heatmap(x, Rowv = NULL, Colv = if(symm)"Rowv" else NULL,  
distfun = dist, hclustfun = hclust,  
reorderfun = function(d, w) reorder(d, w),  
add.expr, symm = FALSE, revC = identical(Colv, "Rowv"),  
scale = c("row", "column", "none"), na.rm = TRUE,  
margins = c(5, 5), ColSideColors, RowSideColors,  
cexRow = 0.2 + 1/log10(nr), cexCol = 0.2 + 1/log10(nc),  
labRow = NULL, labCol = NULL, main = NULL,  
xlab = NULL, ylab = NULL,  
keep.dendro = FALSE, verbose = getOption("verbose"), ...)
```

How to read Heat Map



Summary

Clustering

- Clustering is widely used in bioinformatics
- Clustering can be implemented by using built-in functions of R
- Clustering can be visualized as heat map and dendrogram

Next

- 1 Overview of Machine Learning
- 2 Unsupervised Learning
- 3 **Supervised Learning**
 - Supervised Learning in Bioinformatics
 - Support Vector Machine and its Applications in Bioinformatics
 - Summary
- 4 Future
- 5 Big Data

Supervised Learning: Classification

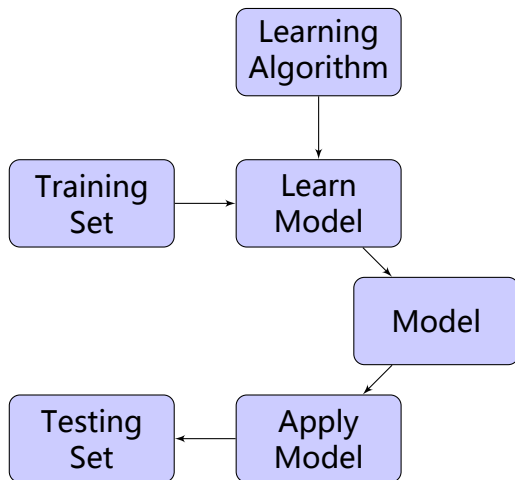
Classification

Assigning objects to one of several predefined categories.

Definition

Classification is the task of learning a target function f that maps each attribute set x to one of the predefined class labels y .

How to solve a classification problem?



General approach for building a classification model

Evaluation

Confusion Matrix for a 2-class problem

	Prediction=1	Prediction=0
Class=1	f_{11}	f_{10}
Class=0	f_{01}	f_{00}

Evaluation

Accuracy and Error Rate

$$\text{Accuracy} = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}}$$

$$= \frac{f_{11} + f_{00}}{f_{11} + f_{10} + f_{00} + f_{01}}$$

$$\text{ErrorRate} = \frac{\text{Number of wrong predictions}}{\text{Total number of predictions}}$$

$$= \frac{f_{10} + f_{01}}{f_{11} + f_{10} + f_{00} + f_{01}}$$

Classification in Bioinformatics

Applications

- Classification of diseases, especially cancer.
- Prediction of clinical outcome.
- Prediction of the function of gene or proteins.
- Prediction of the structure of proteins.
-

Objective

Two Objectives


- ① To build accurate classifiers or predictors
- ② To derive inferences from the results obtained

Challenges

- data inconsistency and missing values
- noise
- normalization
- Dimensionality reduction


SVM

$$K(x_i, x_j) = \phi^T(x_i) \phi(x_j)$$



Kernal

Principle of Support Vector Machines (SVM)



SVM in R

```
library(e1071)
```

SVM

```
library(e1071)
```

```
model = svm(x = exprData[1:400,], y = stages[1:400], cross=5)
```


SVM model

```
model

##
## Call:
## svm.default(x = exprData[1:400, ], y = stages[1:400], cross = 5)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost: 1
##   gamma: 0.001730104
##
## Number of Support Vectors: 200
```

SVM

```
ret = predict(model, exprData[401:500,])  
table(ret, stages[401:500])
```

```
##
```

```
## ret    0    1
```

```
##      0    0    0
```

```
##      1   11   89
```

?

Summary

Classification

- Classification is an important technique for bioinformatics
- SVM is powerful
- Classification algorithm can be implemented in R easily

Next

- 1 Overview of Machine Learning
- 2 Unsupervised Learning
- 3 Supervised Learning
- 4 Future**
- 5 Big Data

Future Reading

- Data Mining for Bioinformatics
- Introduction to Machine Learning
- CRAN Task View: Machine Learning & Statistical Learning: <http://cran.r-project.org/web/views/MachineLearning.html>

Next

- 1 Overview of Machine Learning
- 2 Unsupervised Learning
- 3 Supervised Learning
- 4 Future
- 5 **Big Data**

High Performance Computing

- Task View: High-Performance and Parallel Computing with R

<http://cran.r-project.org/web/views/HighPerformanceComputing.html>

- compiler package and JIT
- Revolution R Enterprise
- Rcpp
- Multi-core
- GPU
- MPI

compiler package

These functions provide an interface to a byte code compiler for R.

```
cmpfun(f, options = NULL) compile(e, env = .GlobalEnv,  
options = NULL) cmpfile(infile, outfile, ascii = FALSE,  
env = .GlobalEnv, verbose = FALSE, options = NULL)  
enableJIT(level)
```

see compiler.R for example

Big Data Framework

- Hadoop:
- Spark: SparkR, AMPLab UC BERKELEY
<http://amplab-extras.github.io/SparkR-pkg/>
- Storm:
 - <https://github.com/allenday/R-Storm>
 - <https://github.com/quintona/storm-r>